

1. Compare the clustering results obtained from the original dataset and PCA-transformed data:

To compare the clustering results, you can visually inspect the clusters formed in both cases and evaluate how similar or different they are. Additionally, you can compare clustering evaluation metrics like silhouette score or Davies–Bouldin index for both cases.

2. Discuss any similarities or differences observed in the clustering results:

After comparing the clustering results, discuss any patterns or differences you observe. Are the clusters formed in both cases similar or do they differ significantly? Are there any clusters present in one representation but not in the other? Analyze the reasons behind these similarities or differences.

3. Reflect on the impact of dimensionality reduction on clustering performance:

Consider how dimensionality reduction using PCA impacted the clustering performance. Did PCA improve or degrade the clustering results compared to using the original dataset? Reflect on how reducing the dimensionality affected the clustering algorithm's ability to capture the underlying structure of the data.

4. Analyze the trade-offs between using PCA and clustering directly on the original dataset:

Analyze the trade-offs between using PCA and clustering directly on the original dataset. Consider factors such as computational efficiency, interpretability of results, and clustering performance. Discuss the advantages and disadvantages of each approach and under what circumstances one might be preferred over the other.

Conclusion:

Summary of Key Findings and Insights:

Provide a brief summary of the key findings from the assignment, including observations from data exploration, clustering results, and comparison between original and PCA-transformed data.

Highlight any significant patterns, similarities, or differences observed during the analysis.

Practical Implications of Using PCA and Clustering in Data Analysis:

Discuss the practical implications of using PCA and clustering in real-world data analysis scenarios.

Explain how PCA helps in reducing the dimensionality of high-dimensional datasets, thereby simplifying the analysis and potentially improving computational efficiency.

Highlight the role of clustering in identifying meaningful patterns or groups within the data, which can be valuable for tasks like customer segmentation, anomaly detection, or recommendation systems.

Recommendations for When to Use Each Technique:

Based on the analysis conducted, provide recommendations for when to use PCA and clustering in data analysis.

Suggest using PCA when dealing with high-dimensional data to reduce dimensionality and remove redundant information while preserving important patterns.

Recommend clustering techniques when there is a need to identify natural groupings or clusters within the data, which can provide valuable insights for decision-making.

Discuss scenarios where using PCA in combination with clustering may yield the best results, such as preprocessing data with PCA before applying clustering algorithms.

pca

May 16, 2024

```
[28]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import warnings
warnings.filterwarnings('ignore')
```

```
[5]: df=pd.read_csv('/content/wine.csv')
```

```
[6]: df.head()
```

```
[6]:
```

| | Type | Alcohol | Malic | Ash | Alcalinity | Magnesium | Phenols | Flavanoids | \ |
|---|------|---------|-------|------|------------|-----------|---------|------------|---|
| 0 | 1 | 14.23 | 1.71 | 2.43 | 15.6 | 127 | 2.80 | 3.06 | |
| 1 | 1 | 13.20 | 1.78 | 2.14 | 11.2 | 100 | 2.65 | 2.76 | |
| 2 | 1 | 13.16 | 2.36 | 2.67 | 18.6 | 101 | 2.80 | 3.24 | |
| 3 | 1 | 14.37 | 1.95 | 2.50 | 16.8 | 113 | 3.85 | 3.49 | |
| 4 | 1 | 13.24 | 2.59 | 2.87 | 21.0 | 118 | 2.80 | 2.69 | |

| | Nonflavanoids | Proanthocyanins | Color | Hue | Dilution | Proline |
|---|---------------|-----------------|-------|------|----------|---------|
| 0 | 0.28 | 2.29 | 5.64 | 1.04 | 3.92 | 1065 |
| 1 | 0.26 | 1.28 | 4.38 | 1.05 | 3.40 | 1050 |
| 2 | 0.30 | 2.81 | 5.68 | 1.03 | 3.17 | 1185 |
| 3 | 0.24 | 2.18 | 7.80 | 0.86 | 3.45 | 1480 |
| 4 | 0.39 | 1.82 | 4.32 | 1.04 | 2.93 | 735 |

```
[7]: df.shape
```

```
[7]: (178, 14)
```

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 178 entries, 0 to 177
Data columns (total 14 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|-----------------|----------------|---------|
| 0 | Type | 178 non-null | int64 |
| 1 | Alcohol | 178 non-null | float64 |
| 2 | Malic | 178 non-null | float64 |
| 3 | Ash | 178 non-null | float64 |
| 4 | Alcalinity | 178 non-null | float64 |
| 5 | Magnesium | 178 non-null | int64 |
| 6 | Phenols | 178 non-null | float64 |
| 7 | Flavanoids | 178 non-null | float64 |
| 8 | Nonflavanoids | 178 non-null | float64 |
| 9 | Proanthocyanins | 178 non-null | float64 |
| 10 | Color | 178 non-null | float64 |
| 11 | Hue | 178 non-null | float64 |
| 12 | Dilution | 178 non-null | float64 |
| 13 | Proline | 178 non-null | int64 |

dtypes: float64(11), int64(3)

memory usage: 19.6 KB

```
[9]: df.describe()
```

```
[9]:
```

| | Type | Alcohol | Malic | Ash | Alcalinity | Magnesium \ |
|-------|------------|------------|------------|------------|------------|-------------|
| count | 178.000000 | 178.000000 | 178.000000 | 178.000000 | 178.000000 | 178.000000 |
| mean | 1.938202 | 13.000618 | 2.336348 | 2.366517 | 19.494944 | 99.741573 |
| std | 0.775035 | 0.811827 | 1.117146 | 0.274344 | 3.339564 | 14.282484 |
| min | 1.000000 | 11.030000 | 0.740000 | 1.360000 | 10.600000 | 70.000000 |
| 25% | 1.000000 | 12.362500 | 1.602500 | 2.210000 | 17.200000 | 88.000000 |
| 50% | 2.000000 | 13.050000 | 1.865000 | 2.360000 | 19.500000 | 98.000000 |
| 75% | 3.000000 | 13.677500 | 3.082500 | 2.557500 | 21.500000 | 107.000000 |
| max | 3.000000 | 14.830000 | 5.800000 | 3.230000 | 30.000000 | 162.000000 |

| | Phenols | Flavanoids | Nonflavanoids | Proanthocyanins | Color \ |
|-------|------------|------------|---------------|-----------------|------------|
| count | 178.000000 | 178.000000 | 178.000000 | 178.000000 | 178.000000 |
| mean | 2.295112 | 2.029270 | 0.361854 | 1.590899 | 5.058090 |
| std | 0.625851 | 0.998859 | 0.124453 | 0.572359 | 2.318286 |
| min | 0.980000 | 0.340000 | 0.130000 | 0.410000 | 1.280000 |
| 25% | 1.742500 | 1.205000 | 0.270000 | 1.250000 | 3.220000 |
| 50% | 2.355000 | 2.135000 | 0.340000 | 1.555000 | 4.690000 |
| 75% | 2.800000 | 2.875000 | 0.437500 | 1.950000 | 6.200000 |
| max | 3.880000 | 5.080000 | 0.660000 | 3.580000 | 13.000000 |

| | Hue | Dilution | Proline |
|-------|------------|------------|------------|
| count | 178.000000 | 178.000000 | 178.000000 |
| mean | 0.957449 | 2.611685 | 746.893258 |
| std | 0.228572 | 0.709990 | 314.907474 |
| min | 0.480000 | 1.270000 | 278.000000 |
| 25% | 0.782500 | 1.937500 | 500.500000 |

| | | | |
|-----|----------|----------|-------------|
| 50% | 0.965000 | 2.780000 | 673.500000 |
| 75% | 1.120000 | 3.170000 | 985.000000 |
| max | 1.710000 | 4.000000 | 1680.000000 |

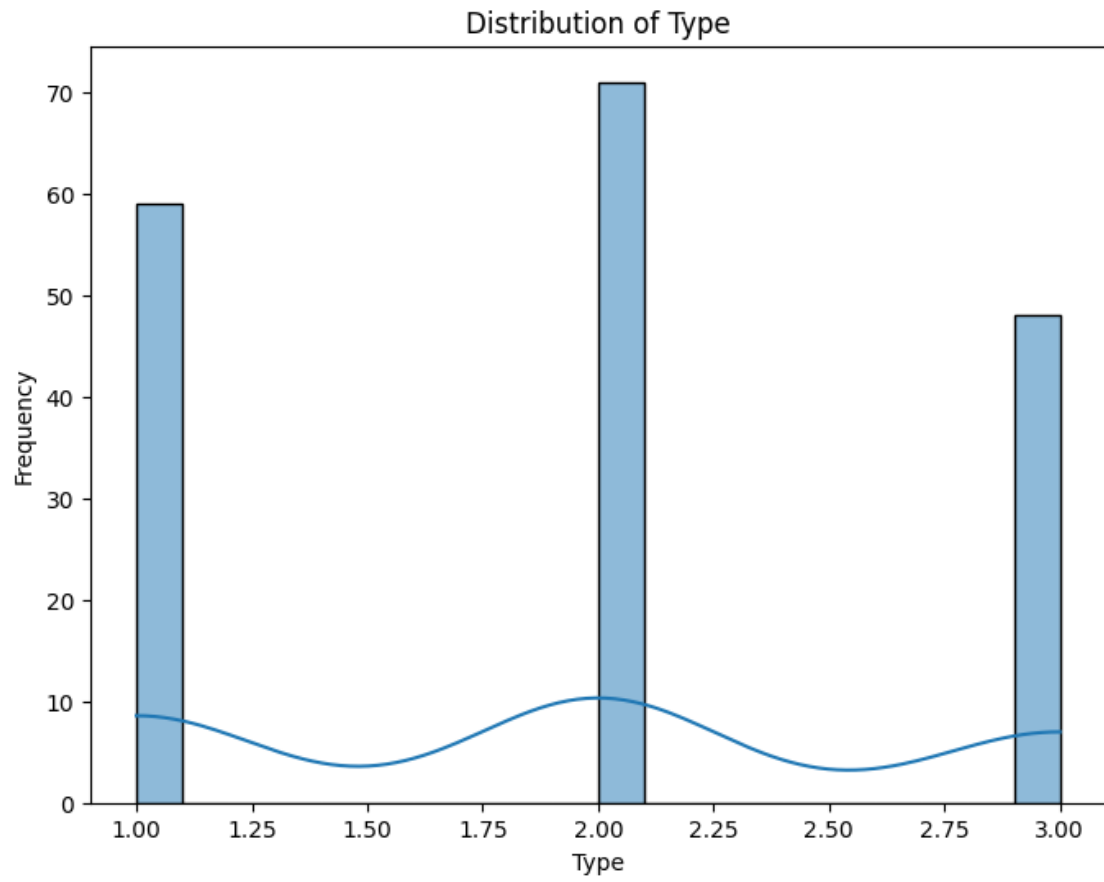
```
[10]: df.isnull().sum()
```


```
[10]: Type      0
      Alcohol    0
      Malic      0
      Ash        0
      Alcalinity  0
      Magnesium  0
      Phenols    0
      Flavanoids  0
      Nonflavanoids  0
      Proanthocyanins  0
      Color      0
      Hue        0
      Dilution  0
      Proline    0
      dtype: int64
```

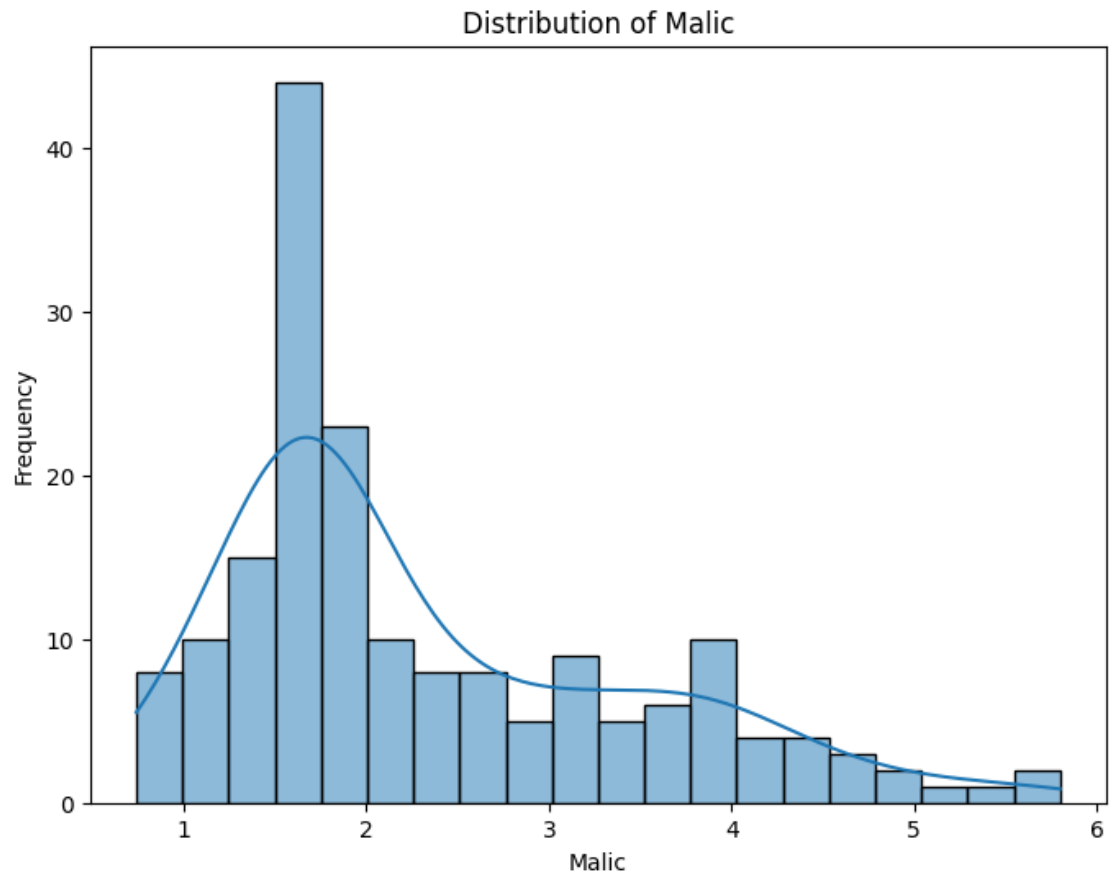
2. Examine the distribution of features using histograms, box plots, or density plots

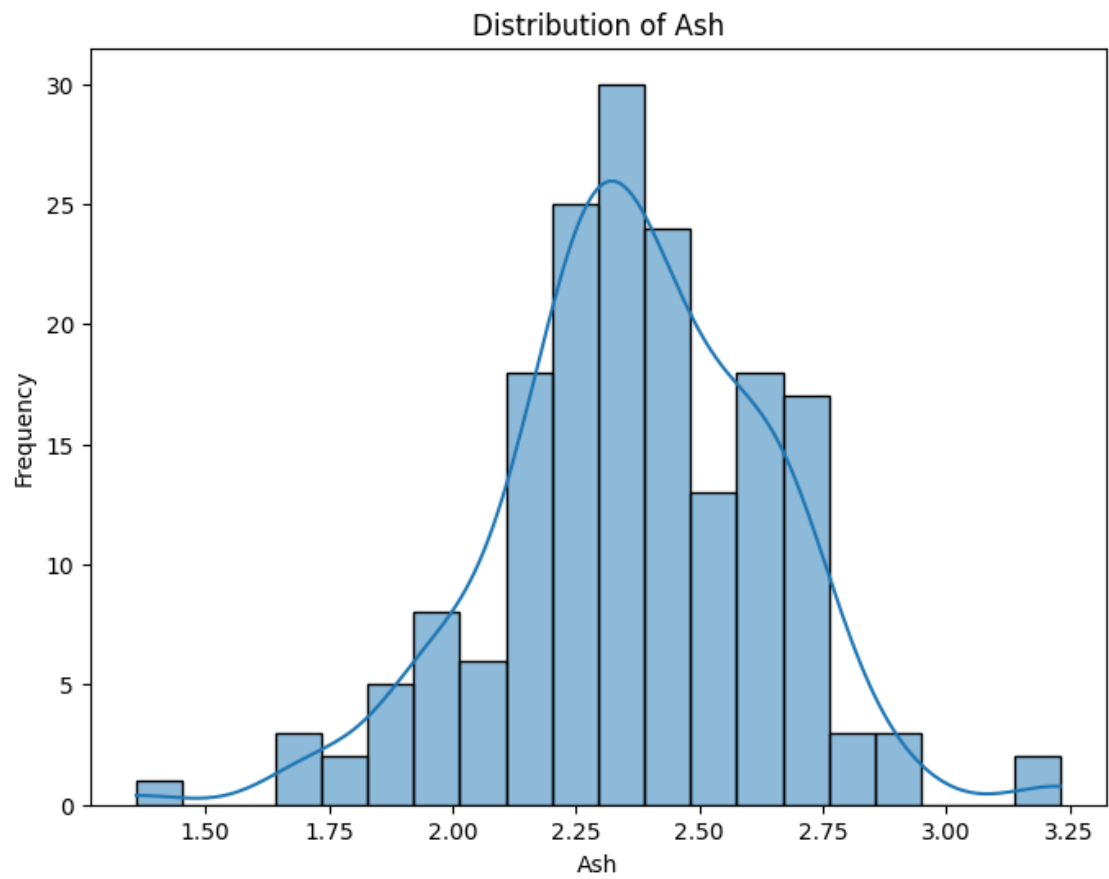
```
[14]: # Examine the distribution of numerical features
numerical_features = df.select_dtypes(include=['float64', 'int64'])
for column in numerical_features.columns:
    plt.figure(figsize=(8, 6))
    sns.histplot(df[column], bins=20, kde=True)
    plt.title(f'Distribution of {column}')
    plt.xlabel(column)
    plt.ylabel('Frequency')
    plt.show()

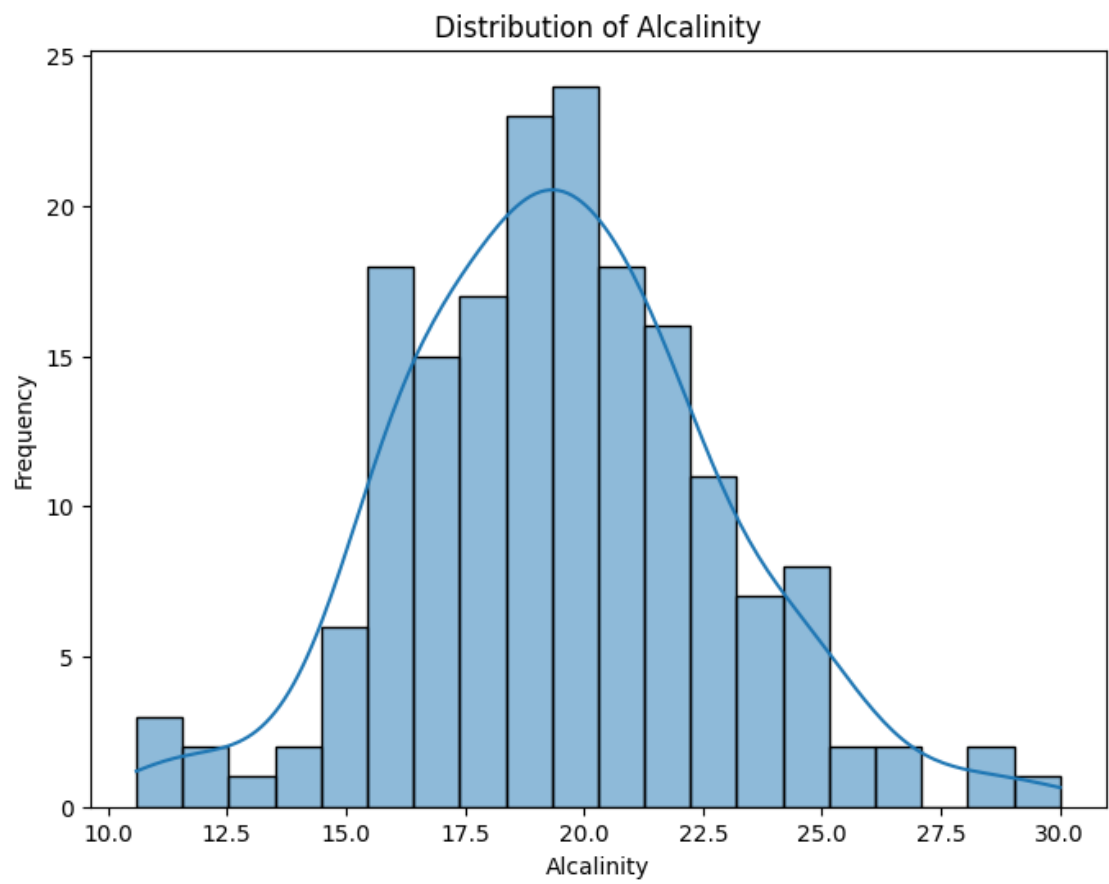
# Examine the distribution of categorical features
categorical_features = df.select_dtypes(include=['object'])
for column in categorical_features.columns:
    plt.figure(figsize=(8, 6))
    sns.countplot(data=df, x=column)
    plt.title(f'Count of {column}')
    plt.xlabel(column)
    plt.ylabel('Count')
    plt.xticks(rotation=45)
    plt.show()
```




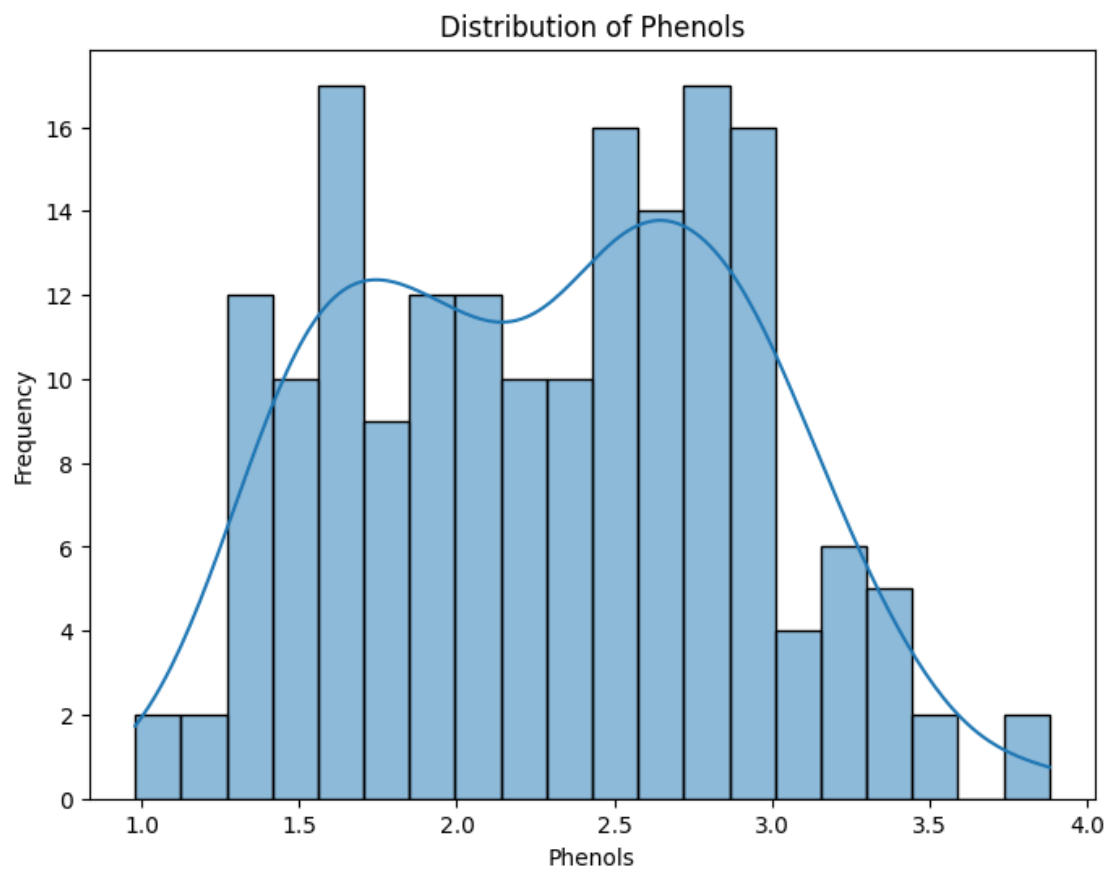


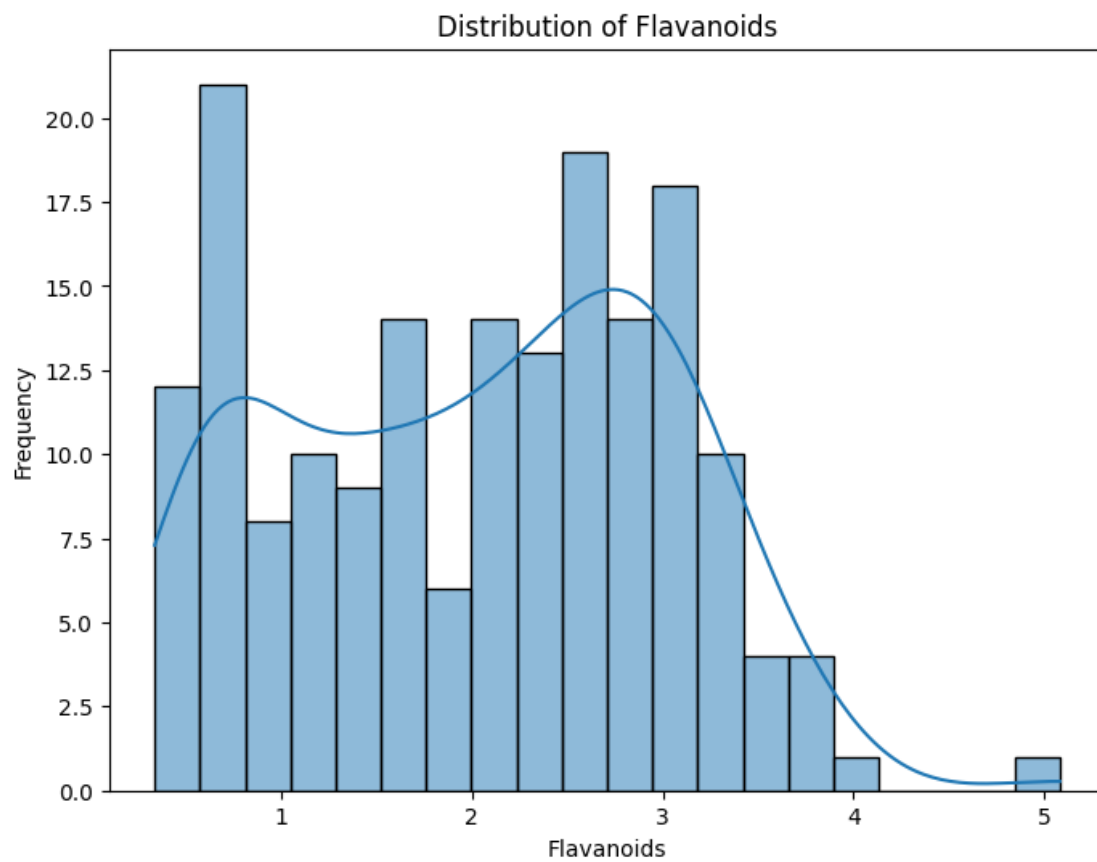


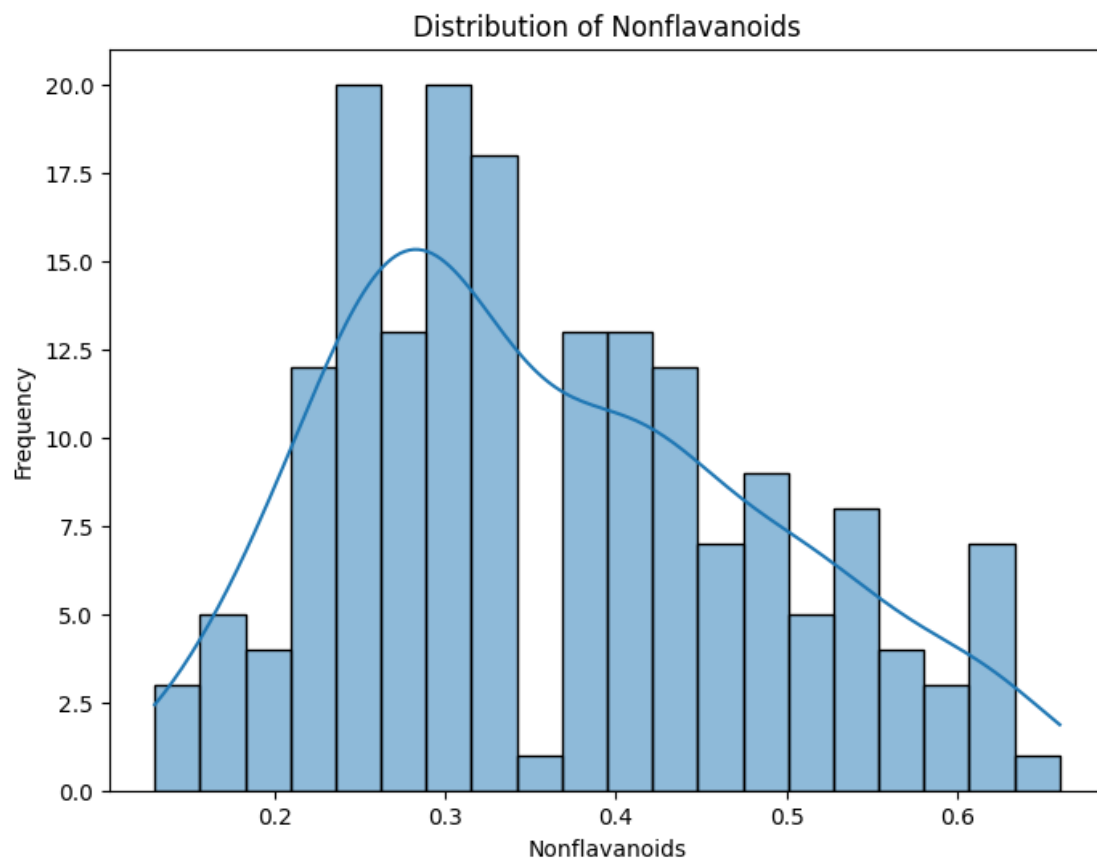





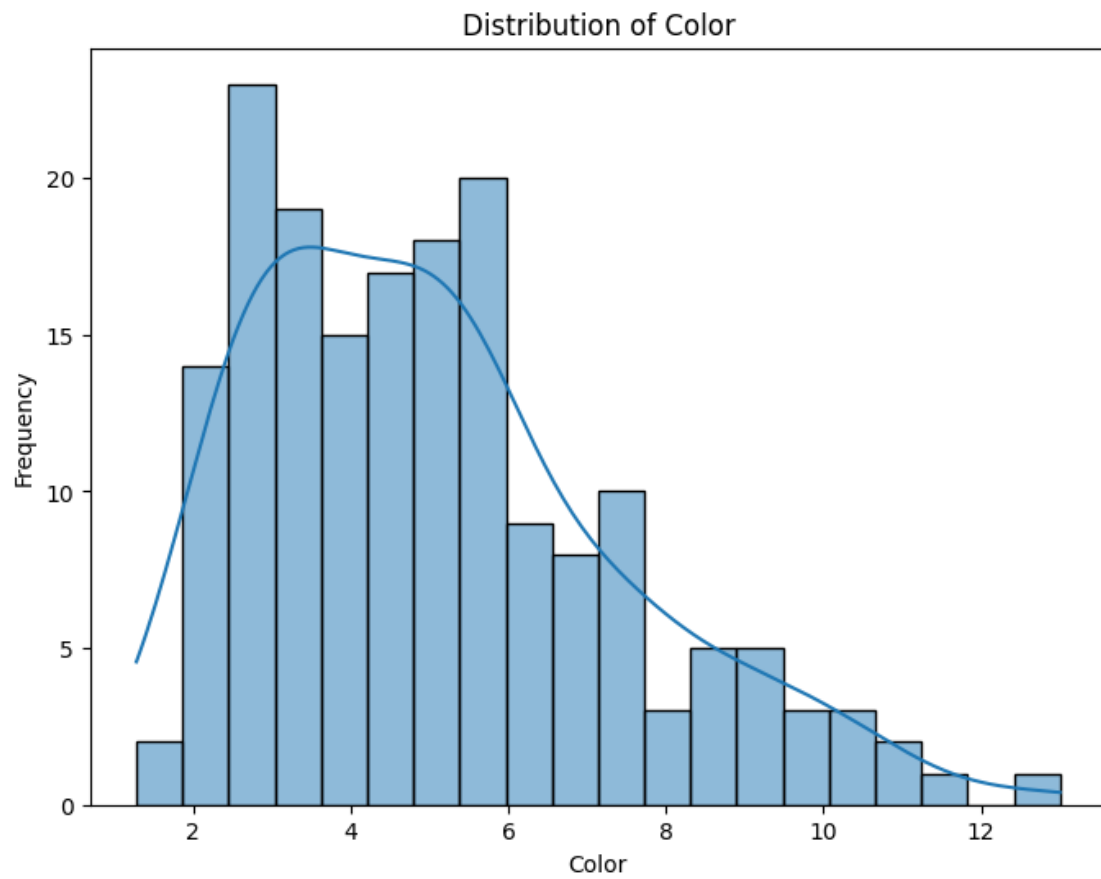


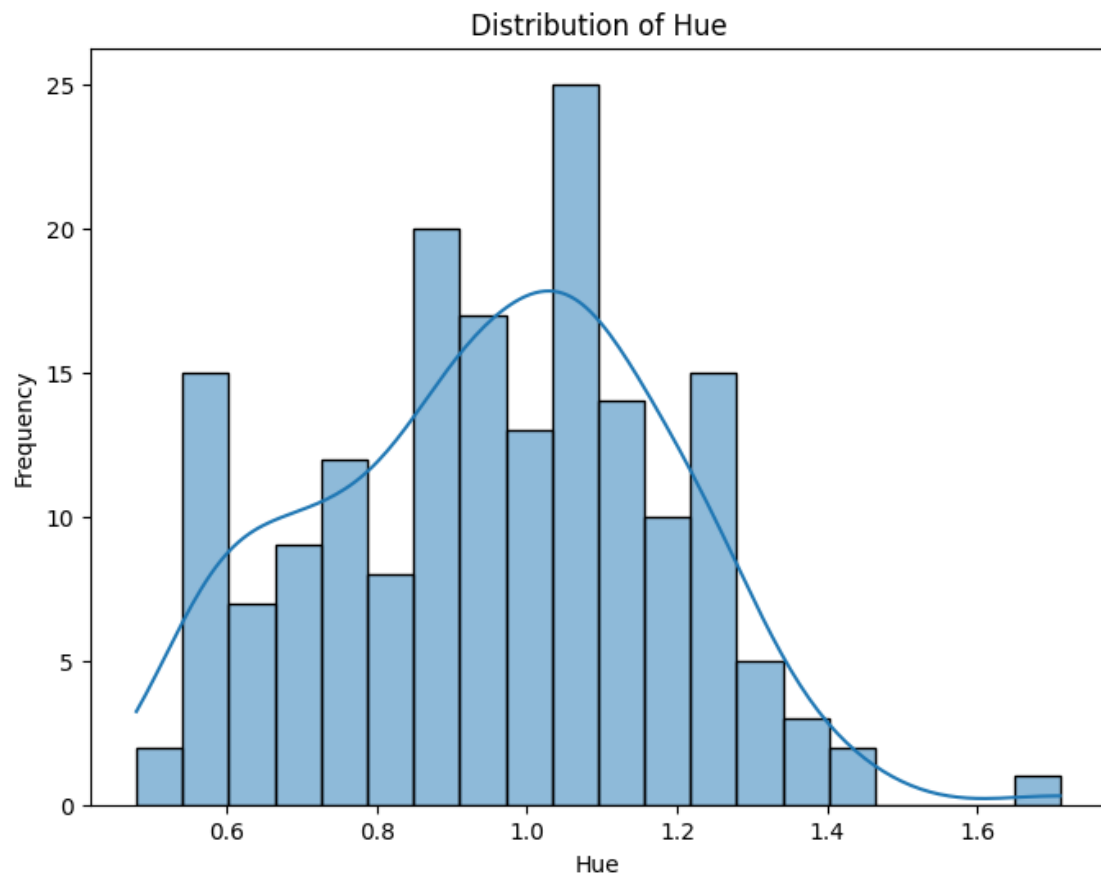


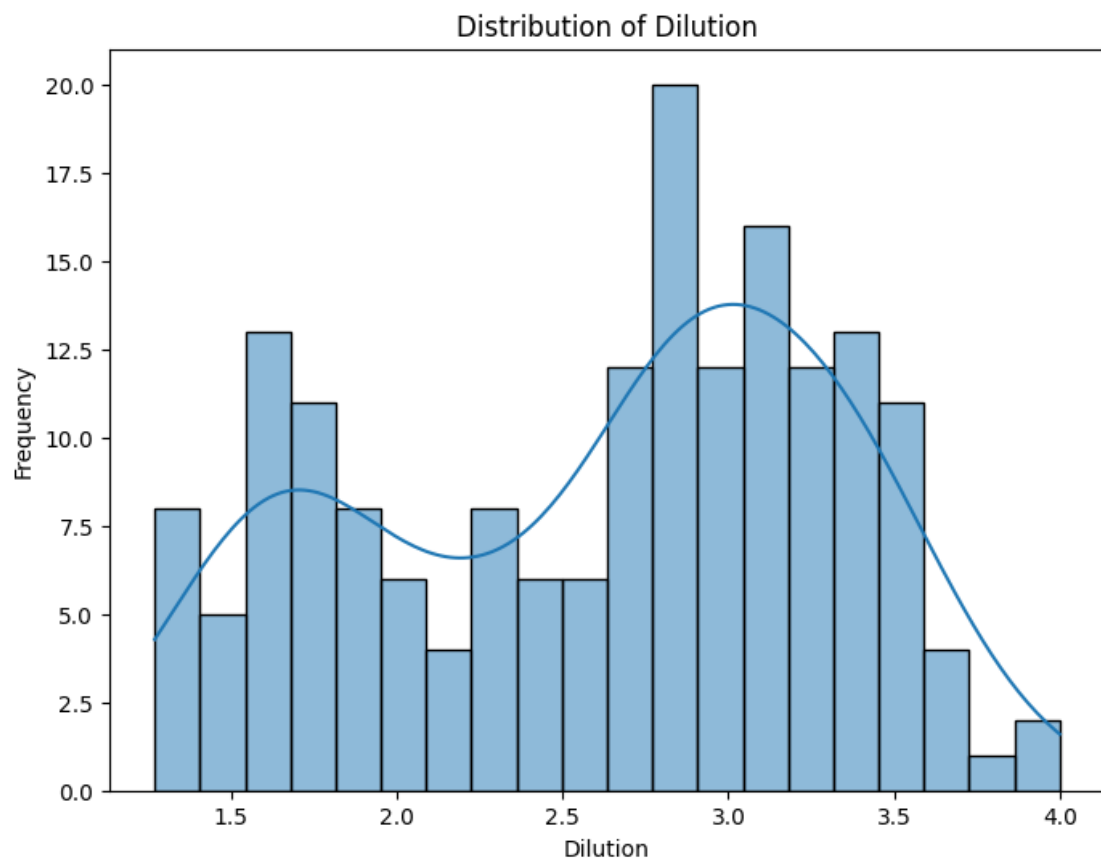


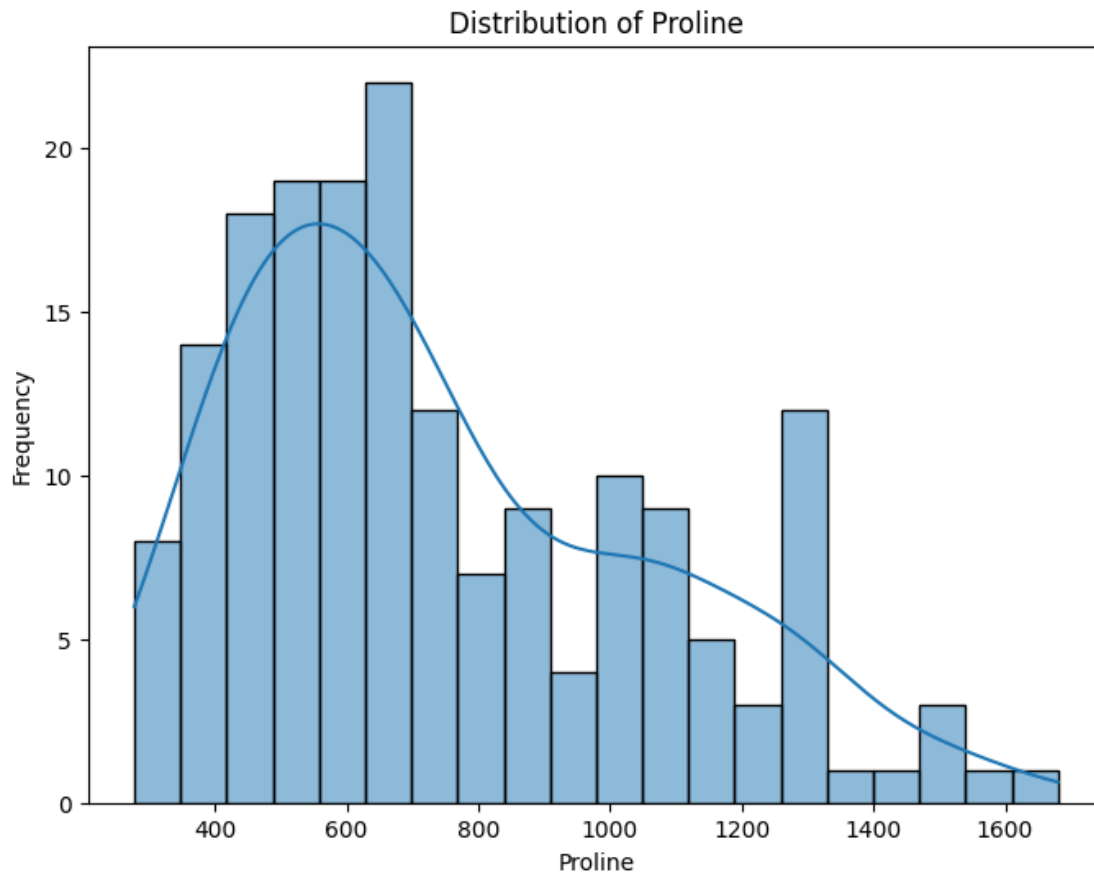








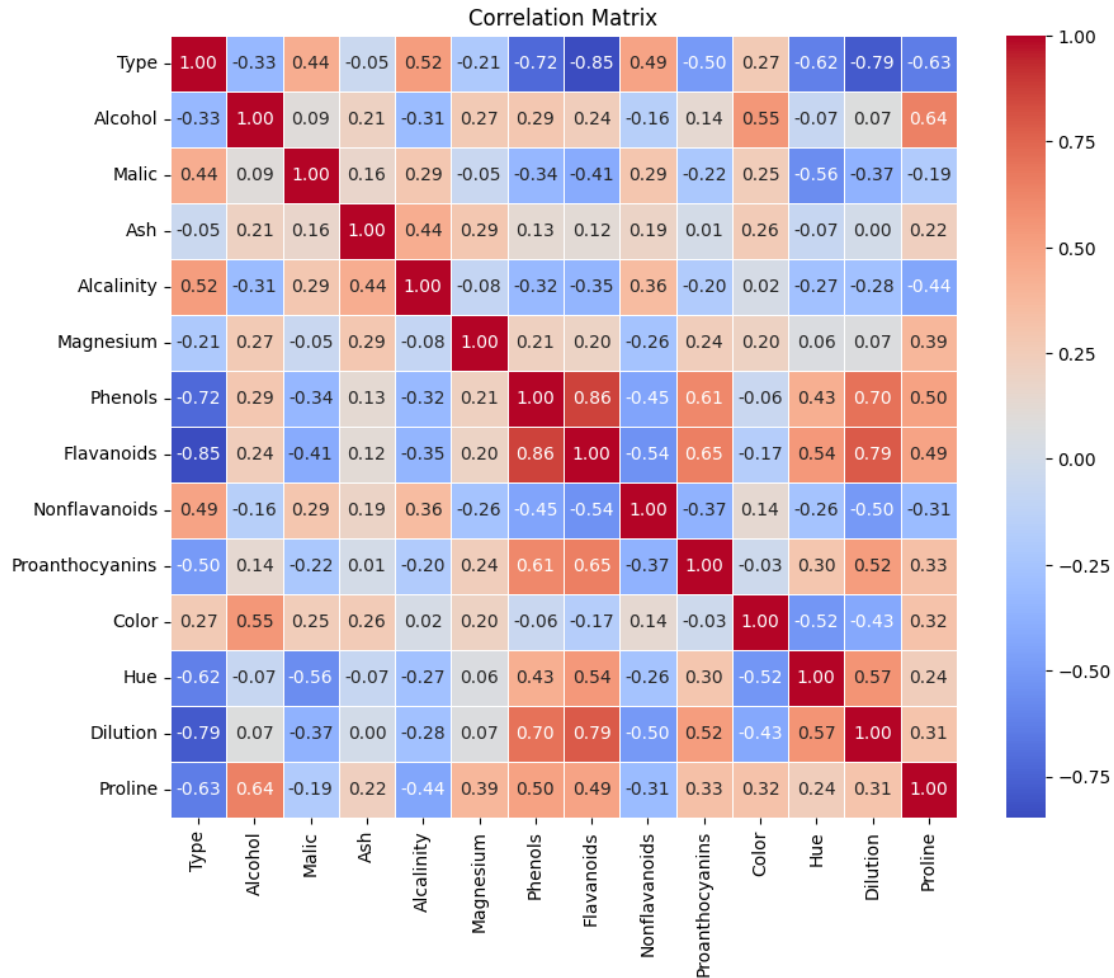




3. Investigate correlations between features

```
[15]: correlation_matrix = df.corr()

# Plot correlation matrix heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
            linewidths=.5)
plt.title('Correlation Matrix')
plt.show()
```



TASK2. 1. Standardize the features

```
[17]: #Standardize the numerical features
scaler = StandardScaler()
standardized_data = scaler.fit_transform(df[numerical_features.columns])
```

2. Initialize PCA with the desired number of components


```
[19]: pca = PCA(n_components=None) # You can specify the number of components or
↳ leave it as None

# Fit PCA to the standardized data
pca.fit(standardized_data)
```

```
[19]: PCA()
```

Determine the optimal number of principal components:

```
[20]: plt.figure(figsize=(10, 6))
plt.plot(range(1, len(pca.explained_variance_ratio_) + 1), pca.
        explained_variance_ratio_, marker='o', linestyle='-')
plt.title('Scree Plot')
plt.xlabel('Number of Components')
plt.ylabel('Explained Variance Ratio')
plt.grid(True)
plt.show()
```



Transform the original dataset into the principal components:

```
[21]: principal_components = pca.transform(standardized_data)
```

TASK 3: Apply a clustering algorithm (e.g., K-means)


```
[25]: num_clusters = 3 # You can adjust this number based on your dataset
kmeans = KMeans(n_clusters=num_clusters, random_state=42)
```

```
[26]: # Fit K-means to the original dataset
kmeans.fit(df[numerical_features.columns])
```

```
[26]: KMeans(n_clusters=3, random_state=42)
```

```
[27]: # Transform the data into principal components for visualization
principal_components_for_vis = pca.transform(standardized_data) # Assuming you
↳ already have fitted PCA

# Plot the clusters in 2D using the first two principal components
plt.figure(figsize=(10, 8))
plt.scatter(principal_components_for_vis[:, 0], principal_components_for_vis[:, 1],
↳ c=kmeans.labels_, cmap='viridis')
plt.title('K-means Clustering (2D)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster')
plt.show()
```



Evaluate the clustering performance

```
[29]: # Calculate the silhouette score
silhouette_avg = silhouette_score(df[numerical_features.columns], kmeans.
    ↪ labels_)
print("Silhouette Score:", silhouette_avg)
```

Silhouette Score: 0.5711220218931753

Task 4: Clustering with PCA Data:

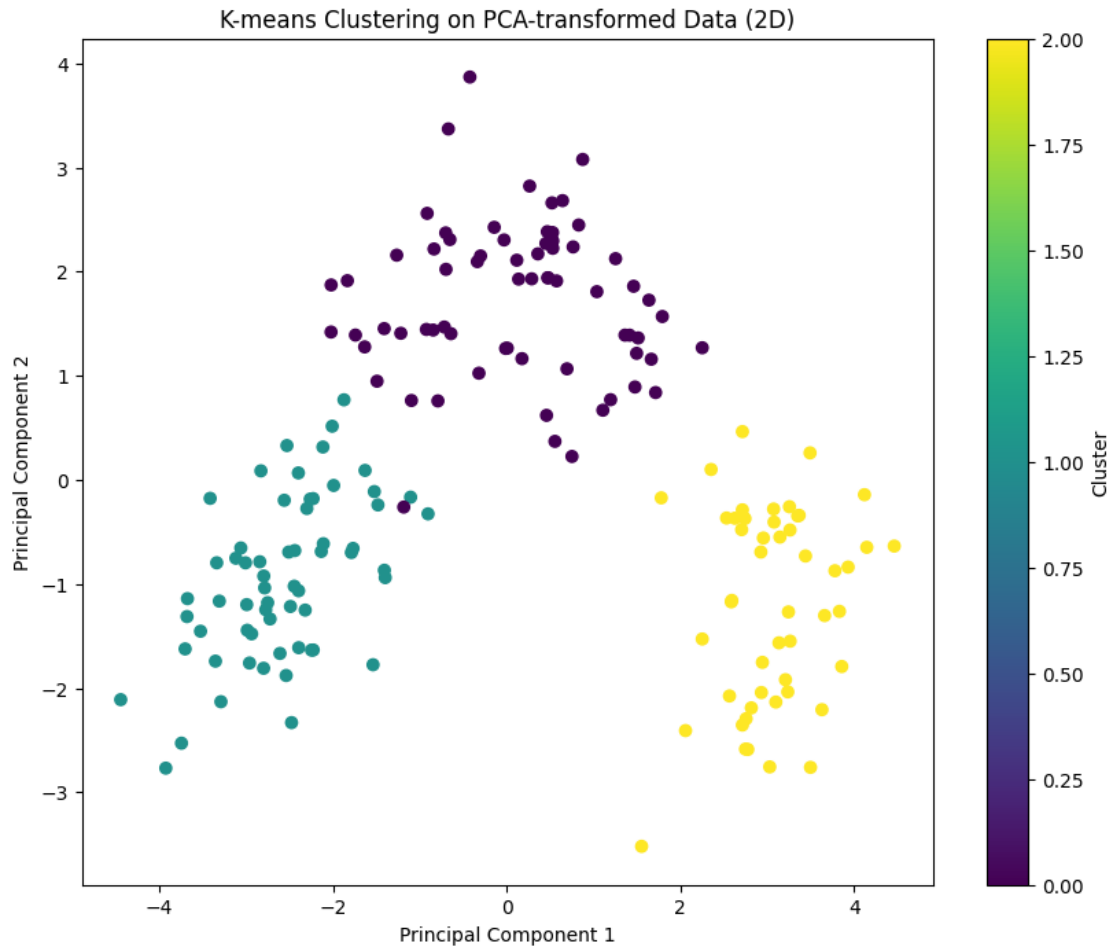
```
[31]: # Initialize K-means with the desired number of clusters
kmeans_pca = KMeans(n_clusters=num_clusters, random_state=42)

# Fit K-means to the PCA-transformed dataset
kmeans_pca.fit(principal_components)
```

```
[31]: KMeans(n_clusters=3, random_state=42)
```

Visualize the clustering results obtained from PCA-transformed data

```
[32]: # Plot the clusters in 2D using the first two principal components
plt.figure(figsize=(10, 8))
plt.scatter(principal_components[:, 0], principal_components[:, 1],
    ↪ c=kmeans_pca.labels_, cmap='viridis')
plt.title('K-means Clustering on PCA-transformed Data (2D)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster')
plt.show()
```



Compare the clustering results from PCA-transformed data with those from the original dataset

```
[33]: # Plot the clusters in 2D using the first two principal components for both
      ↪ datasets
plt.figure(figsize=(18, 8))


plt.subplot(1, 2, 1)
plt.scatter(principal_components_for_vis[:, 0], principal_components_for_vis[:, 1],
            ↪ c=kmeans.labels_, cmap='viridis')
plt.title('K-means Clustering on Original Data (2D)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster')

plt.subplot(1, 2, 2)
plt.scatter(principal_components[:, 0], principal_components[:, 1],
            ↪ c=kmeans_pca.labels_, cmap='viridis')
```



```
plt.title('K-means Clustering on PCA-transformed Data (2D)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.colorbar(label='Cluster')

plt.tight_layout()
plt.show()
```



TASK5